

- Dataset can be found at [Medical Cost Personal Datasets \(https://www.kaggle.com/mirichoi0218/insurance\)](https://www.kaggle.com/mirichoi0218/insurance)
- More about K-Means clustering at [Linear Regression \(https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html)

```
In [1]: ▶ import pandas as pd
import numpy as np

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
In [2]: ▶ df = pd.read_csv("insurance.csv")
df.head()
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

## What does the dataset contain ?

age: age of primary beneficiary

sex: insurance contractor gender, female, male

bmi: Body mass index, providing an understanding of body, weights that are relatively high or low relative to height, objective index of body weight (kg / m ^ 2) using the ratio of height to weight, ideally 18.5 to 24.9

children: Number of children covered by health insurance / Number of dependents

smoker: Smoking

insurance charges

region: the beneficiary's residential area in the US, northeast, southeast, southwest, northwest.

charges: Individual medical costs billed by health insurance

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [4]: `df['sex'].replace({'male': 0, 'female': 1}, inplace=True)`  
`df['smoker'].replace({'yes': 0, 'no': 1}, inplace=True)`  
`df = pd.concat([df, pd.get_dummies(df['region'], prefix='region')], axis=1)`  
`df.drop(columns=['region'], inplace=True)`  
`df.head()`

Out[4]:

	age	sex	bmi	children	smoker	charges	region_northeast	region_northwest	region_southeast	region_southwest
0	19	1	27.900	0	0	16884.92400	0	0	0	1
1	18	0	33.770	1	1	1725.55230	0	0	1	0
2	28	0	33.000	3	1	4449.46200	0	0	1	0
3	33	0	22.705	0	1	21984.47061	0	1	0	0
4	32	0	28.880	0	1	3866.85520	0	1	0	0

## What is the algorithm

Linear regression is a supervised regression algorithm.

Linear regression is a linear approach to modelling the relationship between a scalar response and one or more explanatory variables (also known as dependent and independent variables)

## Advantages and Disadvantages of the algorithm

Advantages:

- \* Linear regression performs exceptionally well for linearly separable data
- \* Easier to implement, interpret and efficient to train
- \* It handles overfitting pretty well using dimensionally reduction techniques, regularization, and cross-validation
- \* One more advantage is the extrapolation beyond a specific data set

Disadvantages:

- \* The assumption of linearity between dependent and independent variables
- \* It is often quite prone to noise and overfitting
- \* Linear regression is quite sensitive to outliers
- \* It is prone to multicollinearity

## How is it performed on the dataset

```
In [5]: ▶ df.head()
```

Out[5]:

	age	sex	bmi	children	smoker	charges	region_northeast	region_northwest	region_southeast	region_southwest
0	19	1	27.900	0	0	16884.92400	0	0	0	1
1	18	0	33.770	1	1	1725.55230	0	0	1	0
2	28	0	33.000	3	1	4449.46200	0	0	1	0
3	33	0	22.705	0	1	21984.47061	0	1	0	0
4	32	0	28.880	0	1	3866.85520	0	1	0	0

```
In [6]: ▶ X = df.drop(columns=['charges'])  
y = df[['charges']]
```

```
In [7]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.3)
```

```
In [8]: ss = StandardScaler()

X_train_std = ss.fit_transform(X_train)
X_test_std = ss.transform(X_test)
```

```
In [9]: regressor = LinearRegression()
regressor.fit(X_train_std, y_train)
```

Out[9]: LinearRegression()

```
In [10]: regressor.score(X_test_std, y_test) ## Accuracy score on test dataset
```

Out[10]: 0.769611805436901

```
In [11]: X_full = ss.transform(X)
predictions = regressor.predict(X_full)
df['Prediction'] = predictions
df.head()
```


Out[11]:

	age	sex	bmi	children	smoker	charges	region_northeast	region_northwest	region_southeast	region_southwest	Pred
0	19	1	27.900	0	0	16884.92400	0	0	0	1	25027.2
1	18	0	33.770	1	1	1725.55230	0	0	1	0	3669.9
2	28	0	33.000	3	1	4449.46200	0	0	1	0	6862.4
3	33	0	22.705	0	1	21984.47061	0	1	0	0	3788.6
4	32	0	28.880	0	1	3866.85520	0	1	0	0	5681.8

```
In [12]: print("The accuracy score of KNN on the dataset is: {}".format(regressor.score(X_full, y)))
```

The accuracy score of KNN on the dataset is: 0.7506272930769431

```
In [13]: ▶ import matplotlib.pyplot as plt
plt.scatter(df['charges'].values, df['Prediction'].values)
plt.xlabel("Actual charges")
plt.ylabel("Predicted charges")
plt.show();
```



## Summary

- The features do not seem to be linearly dependency.
- Linear Regression is limited to linearly dependent data.
- The performance of the model could be increased by feature engineering, mapping the data to a linearly separable space.

```
In [ ]: ▶
```

